

Parallel computation of community structures in graphs

Department of Computer Science
Independent Study by Bobby Srisan
Advisor: Dr. Martin Burtscher
May 1, 2020



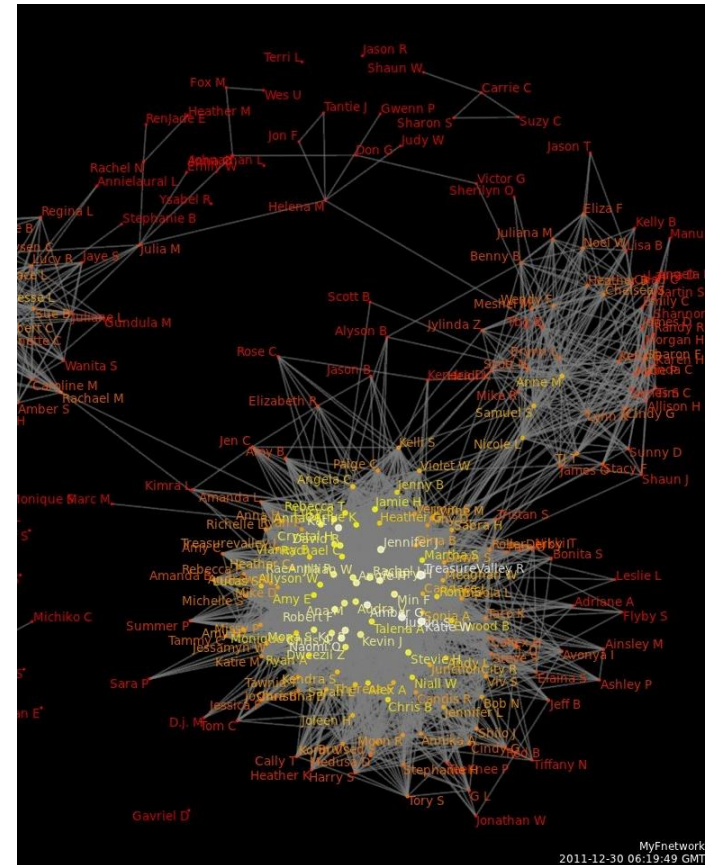
The rising STAR of Texas

Overview

- ❖ Background of communities in network analysis
- ❖ Parallelization project
- ❖ Implementation
- ❖ Results
- ❖ Analysis
- ❖ Summary

Network Analysis

- ❖ Information systems and complex networks are represented as graphs
- ❖ Nodes/vertices connected by links/edges
- ❖ Use in social networks, mobile phone networks, web
- ❖ Organization in the randomness



Data visualization of Facebook relationships by the third-party app MyNetwork
<https://commons.wikimedia.org/wiki/User:Kencf0618>

Communities in Networks

Communities: subunits in a network



Identifying communities helps uncover a-priori unknown ordering

Protein to protein interactions

Communities with shared language dialects

Different group identities discussing a Twitter trend

Community Detection

- ❖ General idea: Optimized network partitioning
 - Communities of densely connected nodes and sparse interconnections between communities
- ❖ Precise formulations are known to be computationally intractable
- ❖ “Goodness” of partition measured by modularity value Q
- ❖ Objective: maximize Q modularity

$$Q = \frac{1}{2m} \sum_{i,j} \left[A_{ij} - \frac{k_i k_j}{2m} \right] \delta(c_i, c_j)$$

where

A_{ij} : weight of link between i and j

k_i : sum of link weights attached to node i


c_i : community of node i

$\delta(c_i, c_j)$: 1 if $c_i == c_j$, else 0

m = sum of link weights in graph

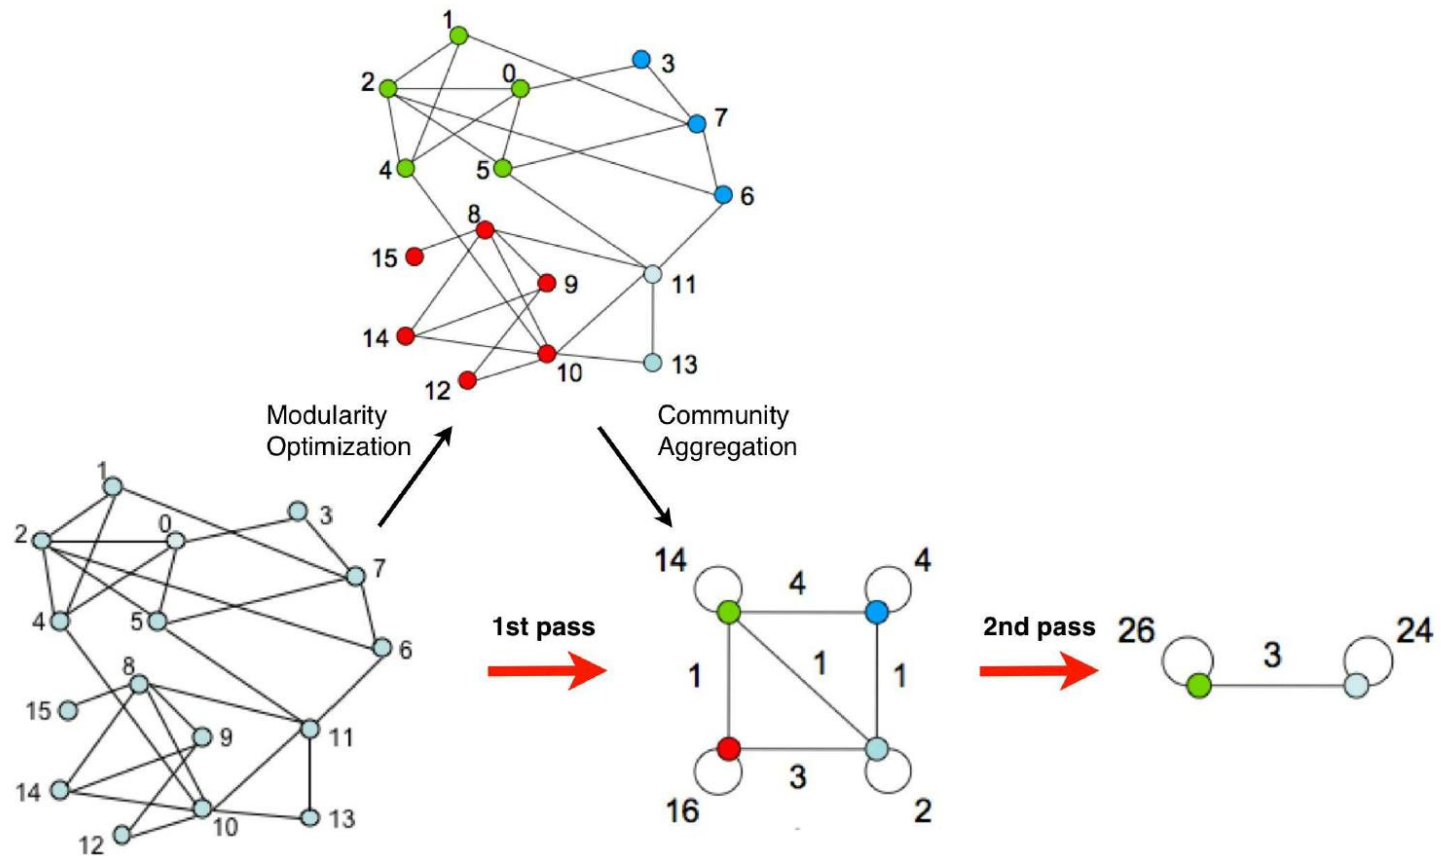
Louvain Method of Community Detection

- ❖ Rapidly finds high modularity partitions of large networks;
Blondel et al.
- ❖ Unfolds hierarchical community structure for the network
- ❖ Unsupervised clustering



[1] Blondel, Vincent D.; Guillaume, Jean-Loup; Lambiotte, Renaud; Lefebvre, Etienne (2008). "Fast unfolding of communities in large networks". Journal of Statistical Mechanics: Theory and Experiment. 2008

Louvain Algorithm



[1] Blondel, Vincent D.; Guillaume, Jean-Loup; Lambiotte, Renaud; Lefebvre, Etienne (2008). "Fast unfolding of communities in large networks". Journal of Statistical Mechanics: Theory and Experiment. 2008

Room for Improvement?

- ❖ While Louvain outperforms other methods, graphs of many millions of nodes/links will take minutes
 - Example: “Web uk-2005” has 39M/783M nodes/links took 738s to achieve 0.979 modularity
- ❖ Explore parallel processing using GPUs to further reduce overall computation time in partitioning nodes into communities

Modularity Optimization/Computation Comparison							
Graph name	Karate	Arxiv	Internet	Web nd.edu	Phone	Web uk-2005	Web WebBase 2001
Nodes/links	34 / 77	9k / 24k	70k / 351k	325k / 1M	2.6M / 6.3M	39M / 783M	118M / 1B
Clauset, Newman, and Moore method	.38 / 0s	.772 / 3.6s	.692 / 799s	.927 / 5034s	-/-	-/-	-/-
Pons and Latapy method	.42 / 0s	.757 / 3.3s	.729 / 575s	.895 / 6666s	-/-	-/-	-/-
Wakita and Tsurumi method	.42 / 0s	.761 / 0.7s	.667 / 62s	.898 / 248s	.56 / 464s	-/-	-/-
Louvain Method	.42 / 0s	.813 / 0s	.781 / 1s	.935 / 3s	.769 / 134s	.979 / 738s	.984 / 152min

[1] Blondel, Vincent D.; Guillaume, Jean-Loup; Lambiotte, Renaud; Lefebvre, Etienne (2008). "Fast unfolding of communities in large networks". Journal of Statistical Mechanics: Theory and Experiment. 2008

* modularity value/run time

** -/- means the method took over 24hrs to run.

Parallelizing the Louvain Method

Porting Strategy



Understand general description of algorithm



Top down read of serial code: Start at general operation and drill into specific functions.



Understand flow and data structures



Identify what is calculated, modified, and accounted



Profile serial code and identify functions where time is spent



Identify parallelizable section(s) where most time can be gained back

Serial Timing

- ❖ Serial code: Most time spent computing “one level”

1741849860 function calls (1740434396 primitive calls) in 2173.463 seconds

Ordered by: cumulative time

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.004	0.004	2173.463	2173.463	{built-in method builtins.exec}
1	0.093	0.093	2173.458	2173.458	<string>:1(<module>)
1	0.175	0.175	2173.365	2173.365	community_louvain.py:161(best_partition)
1	2.714	2.714	2171.409	2171.409	community_louvain.py:253(generate_dendrogram)
6	433.219	72.203	2059.197	343.200	community_louvain.py:463(__one_level)


TEXAS

STATE
UNIVERSITY
The rising STAR of Texas

-

Computing One Level

- ❖ Larger size => more nodes and neighbors to iterate through
- ❖ So, assign GPU threads to work on one node and its neighbors
- ❖ All nodes simultaneously search for their best community
- ❖ Perform a global consensus on which one has the most to gain



Computing Modularity Change

$$\Delta Q = \left[\frac{\Sigma_{in} + k_{i,in}}{2m} - \left(\frac{\Sigma_{tot} + k_i}{2m} \right)^2 \right] - \left[\frac{\Sigma_{in}}{2m} - \left(\frac{\Sigma_{tot}}{2m} \right)^2 - \left(\frac{k_i}{2m} \right)^2 \right]$$

Given node i and neighboring community C ,

Σ_{in} sum of links strictly inside community

$k_{i,in}$ sum of links from i to nodes in C

Σ_{tot} sum of links strictly inside community


k_i sum of links to node i

m sum of all links in entire graph

Computing One Level

Problem: Data dependency of shared community arrays; Q is calculated from these

Solution: Keep arrays as read only; do some personal accounting if community under consideration is current community



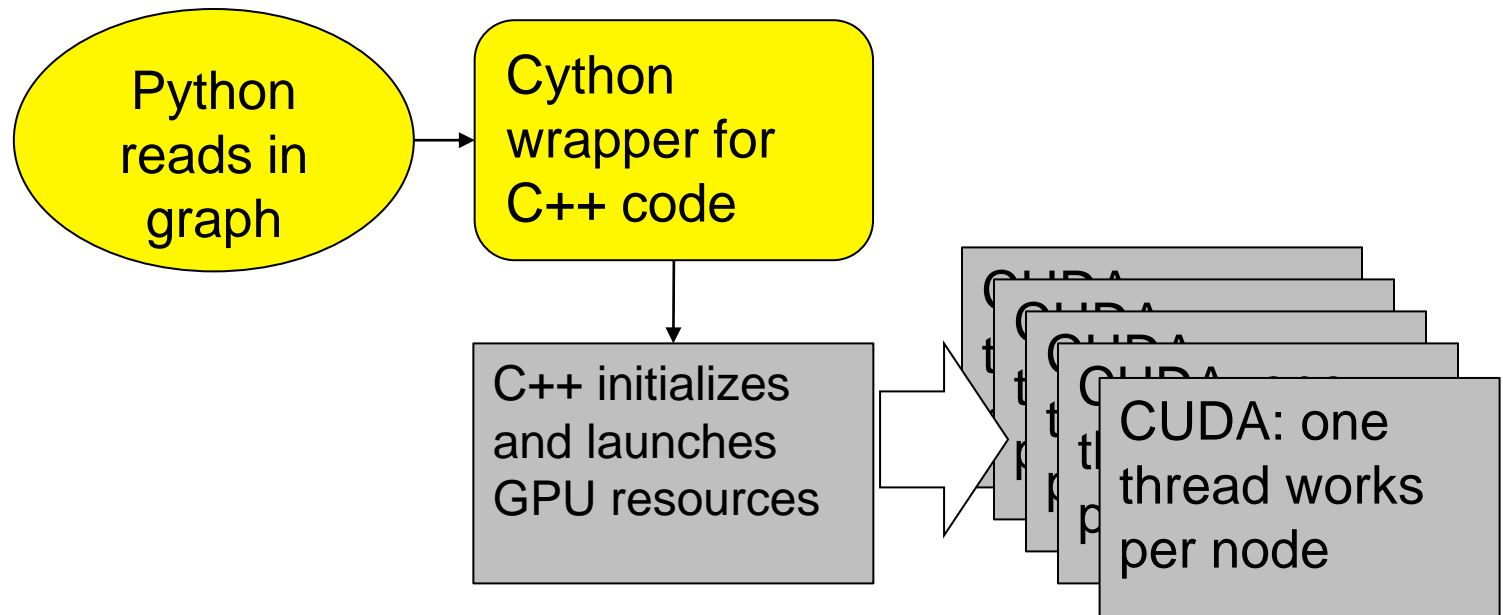
Implementation

Python, Cython, CUDA

- ❖ Python: High-level, general-purpose, “batteries included”
 - ❖ very useful for data analysis when we aren’t too concerned with writing new functions from scratch
- ❖ Cython: Lets us go “under the hood” to optimize performance
- ❖ CUDA C/C++ API: Interface with host (CPU) to GPU device and call operations to take advantage of parallelism
- ❖ During analysis, import the Cython wrapper module to run and get results

Data Passing

- ❖ Read in graph and allocate arrays in CPU and GPU memory
- ❖ CPU initiates the GPU compute kernel
- ❖ GPU's cores run computation for each node in parallel



Hardware

- ❖ CPU: Two Intel Xeon X5690 @ 3.47 GHz
 - Number of Cores: 12 (2x 6-core UMA)
 - Number of Threads: 12
 - Main Memory: 24 GB
- ❖ GPU: NVIDIA Tesla K20 GPU
 - 13 multiprocessors
 - 3.52 TFLOPS single, 1.17 TFLOPS double
 - Threads per multiprocessor: 2048

Data

- ❖ Karate social network
- ❖ Nodes: 34
- ❖ Links: 78
- ❖ Average degree: 5
- ❖ Max degree: 16

Results

Graph	Louvain run time	GPU run time	Louvain modularity	GPU modularity
karate	0.0044	0.0033	0.417	0.394

Average times are in seconds

Serial runs repeated 5 times, randomized nodes

Observations

- ❖ GPU produces a result that original Louvain method can obtain
- ❖ Max of max can easily result in lower modularity at the end
- ❖ GPU initialization overhead

More Parallelism?

- ❖ Even with shared arrays, some information can be read/written without interfering other thread's calculations.
- ❖ Two active nodes are independent if neither writes to same community
 - Process independent nodes in parallel
- ❖ Cautious operations
 - Node A and node B are joining/leaving same community
 - Node A leaving a community that node B is joining, or vice versa

Summary

- ❖ Improve Louvain implementation through speedup
 - ❖ Parallelization of one level
- ❖ Simple global greedy search for highest increase in modularity does not ensure maximum modularity
- ❖ Simultaneously assigning non-adjacent communities
 - ❖ Requires thread safety mechanisms

Thank you!